

Search Problem

Property	BFS	UCS	DFS	DLS	GFS	A*
Completeness	YES, b finite	YES*	NO	NO	NO	YES*
Optimal	NO	YES	NO	NO	NO	YES*
Time	$\mathcal{O}(b^{d+1})$	$\mathcal{O}(b^{1+\lceil C^*/\epsilon \rceil})$	$\mathcal{O}(b^{m+1})$	$\mathcal{O}(b^{\ell+1})$	-	$\mathcal{O}(b^{1+\lceil C^*/\epsilon \rceil})$
Space	$\mathcal{O}(b^{d+1})$	$\mathcal{O}(b^{1+\lceil C^*/\epsilon \rceil})$	$\mathcal{O}(bm)$	$\mathcal{O}(b\ell)$	-	
Frontier	Queue	Priority Queue	Stack	Stack	-	

**Condition:* if b is finite, $cost \geq \epsilon > 0$
Notation: b : max num of successor (branching) of any node (maybe ∞); d : depth of (shallowest) goal node; m : max depth of a node from start node; C^* : optimal cost, $cost \geq \epsilon$
Definition: **Complete:** always find a solution; **Optimal:** find a least-cost solution; **Time C.:** number of nodes generated; **Space C.:** max number of nodes in memory

Uninformed Search

Path Checking: In every path $\langle p_k, c \rangle$, ensure that the final state c is not equal to any ancestors of c along this path: $c \notin \{s_0, s_1, ..., s_k\}$ (make sure does not go back). No increase time and space C.
Cycle Checking: Keep track of all nodes previously expanded during the search using a list (**close list**). When expand n_k to obtain successor c : (1) Ensure c is not equal to any previously expanded node (2) If it is, do not add c to Frontier; Expansive S.C.: $\mathcal{O}(b^{d+1})$
Bread-first Search: children at **end** of Frontier (**queue: last in last out**), extract **first** of the F
Depth-first Search: children at **front** of Frontier (**stack: last in first out**), extract **first** of the F
Depth-limited Search: DFS but only to a pre-specified depth limit D
Iterative Deepening Search: Starting at $d = 0$, loop DLS til solution or fail without cutting off
Uniform Cost Search: expand **least cost** node on F. (**Priority Queue**), same as BFS if all same cost

Informed Search

Greedy Bread-first Search: $f(n) = h(n)$; ignore cost of n ; not complement or optimal
A* Search: $f(n) = g(n) + h(n)$; g : cost path; h : heuristic estimate of cost: run out of time&memory

Poof C. Implies Admissible: $(\forall n_1, n_2, a) \quad h(n_1) \leq C(n_1, a, n_2) + h(n_2) \implies (\forall n) \quad h(n) \leq h^*(n)$
(1)Base case: $k = 1$, one step away from s_g , since consistent: $h(s_i) \geq C(s_i, s_g) + h(s_g)$, since $h(s_g) = 0$, $h(s_i) \geq C(s_i, s_g) = h^*(s_i)$, therefore admissible
(2) Induction step: Suppose assumption holds for every node that is $k - 1$ action away from s_g , given a node s_i , it is k action away from s_g , thus optimal path has $k > 1$ steps
(3) Since h is consistent, have: $h(s_i) \leq C(s_i, s_{i+1}) + h(s_{i+1})$
(4) Note that s_{k+1} is on a least-cost path from s_i , must have the path s_{i+1} to s_g as well, by induction hypothesis have: $h(s_{i+1}) \leq h^*(s_{i+1})$
(5) Combine inequality: $h(s_i) \leq C(s_i, s_{i+1}) + h^*(s_{i+1})$

Proof of Optimal with Consistency: $\forall n_1, n_2, \quad h(n_1) \leq h(n_2) + C(n_1, a, n_2)$
(1) WTS: $\hat{f}_{pop}(s_g) = f(s_i)$, when goal node is pooped, have found optimal
(2) Base case: $\hat{f}_{pop}(s_0) = f(s_0) - h(s_0)$
(3) Induction step: Assume $\forall s_0, s_1, ..., s_k, \hat{f}_{pop}(s_i) = f(s_i)$, given:
$$\hat{f}_{pop}(s_{k+1}) = \hat{g}_{pop}(s_{k+1}) + h(s_{k+1}) \geq g(s_{k+1}) + h(s_{k+1}) = f(s_{k+1})$$

For s_{k+1} is only explored after s_k , require $f(s_i) \leq f(s_{k+1})$, need of consistency of h , pooping s_k
$$\begin{aligned} \hat{f}_{pop}(s_{k+1}) &= \min\{\hat{f}(s_{k+1}), \hat{g}_{pop}(s_k) + c(s_k, s_{k+1}) + h(s_{k+1})\} \\ &\leq \hat{g}_{pop}(s_k) + c(s_k, s_{k+1}) + h(s_{k+1}) \\ &= g(s_k) + c(s_k, s_{k+1}) + h(s_{k+1}) \\ &= g(s_{k+1}) + h(s_{k+1}) \\ &= f(s_{k+1}) \end{aligned}$$

from IH

IDA* Search: reduce memory requirements of A*; cutoff is the f-value rather than the depth; at each iteration, the cutoff is the **smallest f-value** of any node that exceeded the cutoff on the previous iteration; avoids overhead with keeping a sorted queue of nodes, the Frontier occupies **linear space**.

CSPs

- (1) A set of **variables** $V_1, ..., V_n$;
 - (2) A (finite) **domain** of possible values $Dom[V_i]$ for each variable V_i
 - (3) A set of **constraints** $C_1, ..., C_m$,
Unary: over one variable: $C(X) : X = 2$
Binary: over two variable: $C(X, Y) : X + Y \geq 2$
Higher-order: over $i \geq 3$ variable: $All - Diff(V_1, ..., V_n) : V_1 \neq V_2, ..., V_2 \neq V_1, ... V_n \neq V_{n-1}$
 - (4) Each variable V_i can be assigned any value from its domain: $V_i = d$ where $d \in Dom[V_i]$
 - (5) Each constraint C Has a set of variables it operates over, called its **scope**.
 - (6) Solution to a CSP: An assignment of a value to all of the variables such that every constraint is **satisfied**; unsatisfiable if no solution exists.
- Back Tracking Search:** searching through the space of partial assignments, rather than paths. Decide on a suitable value for one variable at a time. Order in which we assign the variables does not matter. If a constraint is falsified during the process of partial assignment, immediately reject the current partial assignment.
Back Tracking Search with Inference: every time assign a value to variable V , check **all constrains over V** and prune values from the current domain of the **unassigned variables** of the constrains
- (1) *Value Assignment:* define current domain (CurDom) of a value; first step to infer other values
 - (2) **Degree Heuristic:** select the variable that is involved in the largest number of constrains on other **unassigned** variables
 - (3) **Minimum Remaining Values Heuristics:** always branch on a variable with the **smallest remaining values** (smallest CurDom)
 - (4) **Least Construing Value Heuristic:** always pick a value in CurDom that **rules out the least domain values** of other neighboring values in the constraint

Games

Properties: two player; **finite** number of states and moves (large- heuristic cutoffs); deterministic (perfect info/observable); **zero-sum**: **fully competitive**, total payoff to all players is constant

- 2 players **Max** and **Min**
- A set of positions P (states of the game)
- A starting positions $P \in P$ (game begins)
- A set of Terminal positions $T \subseteq P$ (game end)
- A set of directed edges E_{Max} between some positions, representing **Max's move**
- A set of directed edges E_{Min} between some positions, representing **Min's move**
- A utility/payoff function $U : T \rightarrow \mathbb{R}$, representing quality each terminal state is for player **Max**

Minmax Search

Max plays a move to change the state to the **highest valued child** $U(S_0) = \max \{U(S_i), ..., U(S_n)\}$
Min plays a move to change the state to the **lowest valued child** $U(S_0) = \min \{U(S_i), ..., U(S_n)\}$
Use **DFS** to save space (finite depth); T.C.: $\mathcal{O}(b^d)$; S.C: $\mathcal{O}(bd)$

Alpha-Beta Pruning

At a Max node s
(1.1) α_s : the **highest value** of s **children** examine so far (changes as examine more children)
(1.2) β : the **lowest value** of s **parent** examine so far (fixed)
(2) If α_s becomes $\geq \beta$, stop expanding children of s ; **Min** never choose to move from s parent, would choose one of s lower valued siblings
At a Min node s
(1.1) α : the **highest value** found so far by s **parent** by previous explored siblings (fixed)
(1.2) β_s : the **lowest value** of value of s **children** examine so far (changes as explore more children)
(2) If α_s becomes $\geq \beta_s$, stop expanding children of s ; **Max** never choose to move from s parent, would choose one of s higher valued siblings

- Set initial values: $\alpha = -\infty$ and $\beta = \infty$
- While backing the utility values up the tree, identity α, β for each node (α/β : best already explored along the path to the root of **MAX/MIN**)
- At every node s , if $\alpha \geq \beta$ **prune** (remaining) children of s (α/β -cuts: pruning of **MAX/MIN** nodes)

Ordering Moves: **Max** prune best if best move for Max explored first; **Min** prune best if best move for Min explored first; can use heuristics to estimate and choose
Effectiveness: no pruning ($\mathcal{O}(b^d)$); if move **ordering is optimal** ($\mathcal{O}(b^{d/2})$)

Bayesian Networks

Probability

\cap : **OR**; \cup : **AND**
Basic Rules: $P(\mathcal{U}) = 1$, $P(A) \in [0, 1]$, $P(A \cup B) = P(A) + P(B) - P(A \cap B)$
Summing out Rule: $P(A) = \sum_{C_i} P(A \cap C_i)$, $P(A | B) = \sum_{C_i} P(A | B \cap C_i)P(C_i | B)$
Normalizing: dividing each number by the sum of the numbers:
(1) normalize $[x_1, x_2, ..., x_k] = [\frac{x_1}{\alpha}, \frac{x_2}{\alpha}, ..., x_k/\alpha]$, where α is the sum of all x_k
(2) normalize $[x_1, x_2, ..., x_k] = [x_1 \cdot \beta, x_2 \cdot \beta, ..., x_k \cdot \beta]$, where β is any constant
Conditional Probability:

$$P(A | B) = \frac{P(A \cap B)}{P(B)}$$

Bayes Rule:

$$P(A | B) = \frac{P(B | A)P(A)}{P(B)}$$

Chain Rule:

$$P(A_1 \cap A_2 \dots \cap A_n) = P(A_n | A_1 \dots \cap A_{n-1}) \cdot P(A_{n-1} | A_1 \dots \cap A_{n-1}) \cdot \dots \cdot P(A_2 | A_1) \cdot P(A_1)$$

Independence: $P(V_1 | V_2) = P(V_1)$ (V_1, V_2 are independent)

$$P(A | B) = P(A) \quad P(A \cap B) = P(A) \cdot P(B)$$

Conditional Independence: A is conditionally independent of B given C

$$P(A | B \cap C) = P(A | C) \quad P(A \cap B | C) = P(A | C) \cdot P(B | C)$$

Joint Distribution: $Dom[V_1] = Dom[V_2] = \{1, 2, 3\}$, $P(V_1, V_2)$ are vector of 9

$$P(v_1 = 1, V_2 = 1), P(V_1 = 1, V_2 = 2), ..., P(V_1 = 2, V_2 = 1), ..., P(V_1 = 3, V_2 = 3)$$

Conditional Probabilty Table (CPT): $Dom[V_1] = Dom[V_2] = Dom[V_3] = \{1, 2, 3\}$, $P(V_1 | V_2, V_3)$ are 27 values; $P(V_1 = 1 | V_2 = 1, V_3 = 1), ..., P(V_1 = 3 | V_2 = 3, V_3 = 3)$

Bayesian Network

Conditional Independence:

$E \rightarrow C \rightarrow A \rightarrow B \rightarrow H$:
(1) B is independent of E , and C , **given** A , A is independent of E **given** C
(2) Computation: $P(H | B, \{A, C, E\}) = P(H | B)$
(3) Chain rule: $P(H, B, A, C, E) = P(H | B, A, C, E)P(B | A, C, E)P(A | C, E)P(C | E)P(E)$
(4) Independence assumption: $P(H, B, A, C, E) = P(H | B)P(B | A)P(A | C)P(C | E)P(E)$
(5) Joint distributions:
 $P(\mathcal{A} = \sum_{x_i \in Dom(X)} P(\mathcal{A} | x_i)P(x_i) = \sum_{x_i \in Dom(X)} P(\mathcal{A} | x_1) \sum_{y_i \in Dom(Y)} P(x_i | y_i)P(y_i)$
Ex. $P(c) = P(c | e)P(e) + P(c | \neg e)P(\neg e)$

Network and Chain Rule

Variable Elimination & Factoring

VE **sum out** the innermost variable, computing a new **function** over **variables in that sum**.

D-Separation (Independence)

Independence: every X_i is conditionally independent of all of its **nondescendants** given its parents

(1) A set of variables E **d-separates** X, Y if it **blocks every undirected path** in the BN between X, Y

Let P be an **undirected path** from X, Y in a BN; let E (evidence) be a set of variables

E **blocks** path P iff there is some node Z on path P such that:

- $Z \in E$ and one arc on P enters (goes into) Z and one leaves (goes out of) Z
- $Z \in E$ and both arcs on P leave Z
- Both arcs on P enter Z and neither Z , nor any of its descendants are in E

(2) X, Y are **conditionally independent** given evidence E if E d-separates X, Y

Knowledge Representation

Representation: **Symbolic** encoding of propositional believed

Reasoning: **Manipulation** of symbolic encoding of propositions to produce propositions that believed by the agent but are **not explicitly stated**

First-order Logic

Syntax: A grammar specifying what are legal syntactic constructs of the representation.

- Propositional variable: **True** or **False** variables
- $A \wedge B$ (conjunction); $A \vee B$ (disjunction); $A \implies B$ (implication); $A \iff B$ (bi-implication)
- A st V of variables; A set of F of function symbols; A set of P of predicate/relation symbols

Let \mathcal{L} be a vocabulary, the set of first-order \mathcal{L} -formulas as defined:

- Atomic formula: $P(t_1, t_2, \dots, t_n)$ where P is an n -ary predicate symbol in \mathcal{L} , and t_n are \mathcal{L} terms
- Negation: $\neg f$, where f is a \mathcal{L} -formula
- Conjunction: $f_1 \wedge f_2 \wedge \dots \wedge f_n$, where f_1, \dots, f_n are \mathcal{L} -formula
- Disjunction $f_1 \vee f_2 \vee \dots \vee f_n$, where f_1, \dots, f_n are \mathcal{L} -formula
- Implication: $f_1 \implies f_2$, where f_1, f_2 are \mathcal{L} -formula
- Existential: $\exists x f$, where x is a variable and f is a \mathcal{L} -formula
- Universal $\forall x f$, where x is a variable and f is a \mathcal{L} -formula

Ex. $AC(x)$: x belongs to Alpine Club; $L(x, y)$: x likes y

Semantics: A formal mapping from syntactic constructs to set theoretic assertions

Truth Assignment: a function τ from the propositional variables into the set of $\{T, F\}$

Let τ be a t.a. extension $\bar{\tau}$ of τ assigns either T, F to every formula and is defined as:

- If $A = x$, where x is a variable, then $\bar{\tau} = \tau(x)$
- $\bar{\tau}(\neg A) = T$, IFF $\bar{\tau}(A) = F$
- $\bar{\tau}(A \wedge B) = T$ IFF $\bar{\tau}(A) = T$ AND $\bar{\tau}(B) = T$
- $\bar{\tau}(A \vee B)$ IFF $\bar{\tau}(A) = T$ OR $\bar{\tau}(B) = T$
- $\bar{\tau}(A \implies B) = F$ IFF $\bar{\tau}(A) = T$ AND $\bar{\tau}(B) = F$

τ **satisfies a set** Φ of formulas IFF τ satisfies all formula in Φ

a formula A is a **logical consequence** of $\Phi \models A$ IFF for every t.a. τ satisfies Φ , then τ satisfies A

Structure: let \mathcal{L} be a first order vocabulary, an \mathcal{L} -structure \mathcal{M} consists:

- Nonempty set M called the **universe (domain) of discourse**
- For each n -ary function symbol $f \in \mathcal{L}$, and associated function $f^{\mathcal{M}} : M^n \rightarrow M$ (if $n = 0$, then f is a constant symbol and $f^{\mathcal{M}}$ is simply an element of M . $f^{\mathcal{M}}$ is called the **extension** of the function symbol f in \mathcal{M})
- For each n -ary predicate symbol $P \in \mathcal{M}$, an assorted relation $P^{\mathcal{M}} \subseteq M^n$. $P^{\mathcal{M}}$ is called the extension of the predicate symbol P in \mathcal{M}

Variable Assignments: let \mathcal{M} be a structure and X be a set of variables. An **object assignment** σ for \mathcal{M} is **mapping** from variables in X to the universe of \mathcal{M}

Recursive definition: let \mathcal{L} be a set of function and predicate symbols

(1) Every variable x is a term;

(2) if f is an n -ary function symbol in \mathcal{L} and t_1, t_2, \dots, t_n are \mathcal{L} -terms, then $f(t_1, t_2, \dots, t_n)$ is a \mathcal{L} -term

Let \mathcal{L} be a vocabulary and \mathcal{M} be an \mathcal{L} -structure, the extension $\bar{\sigma}$ of σ is defined recursively:

(1) For every variable x , $\bar{\sigma}(x) = \sigma(x)$;

(2) For every function symbol $f \in \mathcal{L}$, $\bar{\sigma}(f(t_1, \dots, t_n)) = f^{\mathcal{M}}(\bar{\sigma}(t_1), \dots, \bar{\sigma}(t_n))$

Model Interpretation

For an \mathcal{L} -formula C , $\mathcal{M} \models C[\sigma]$ (\mathcal{M} **satisfies** C under σ , or \mathcal{M} is a **model** of C under σ) is defined recursively on the structure of C as:

- $\mathcal{M} \models P(t_1, \dots, t_n)[\sigma] \iff \langle \bar{\sigma}(t_1), \dots, \bar{\sigma}(t_n) \rangle \in P^{\mathcal{M}}$
- $\mathcal{M} \models (s = t)[\sigma] \iff \bar{\sigma}(s) = \bar{\sigma}(t)$
- $\mathcal{M} \models \neg A[\sigma] \iff \mathcal{M} \not\models A[\sigma]$
- $\mathcal{M} \models (A \vee B)[\sigma] \iff \mathcal{M} \models A[\sigma] \vee \mathcal{M} \models B[\sigma]$
- $\mathcal{M} \models (A \wedge B)[\sigma] \iff \mathcal{M} \models A[\sigma] \wedge \mathcal{M} \models B[\sigma]$
- $\mathcal{M} \models (\forall x A)[\sigma] \iff \mathcal{M} \models A[\sigma(m/x)]$ for all $m \in M$
- $\mathcal{M} \models (\exists x A)[\sigma] \iff \mathcal{M} \models A[\sigma(m/x)]$ for some $m \in M$

$\sigma(m/x)$ is an o.a. function exactly like σ , but maps the variable x to the individual $m \in M$:

(1) For $y \neq x$: $\sigma(m/x)(y) = \sigma(y)$ (2) For x : $\sigma(m/x)(x) = m$

Bounded: an occurrence of $x \in A$ is bounded iff it is an sub-formula of A of the form $\forall x B$ or $\exists x B$; otherwise the occurrence is **free**

In a structure \mathcal{M} , formula with free variables might be true for some object assignments to the free variable and false to others

Sentence: a formula A is closed if it contains no free occurrence of a variable

If σ and σ' agree on the free variables of A , then $\mathcal{M} \models A[\sigma] \iff \mathcal{M} \models A[\sigma']$

Corollary: if A is a **sentence**, then for any object assignments σ , and σ' :

$$\mathcal{M} \models A[\sigma] \iff \mathcal{M} \models A[\sigma']$$

so if A is a sentence (no free variables), σ is irrelevant and omit mention of σ , $\mathcal{M} \models A$

Satisfiability: let Φ be a set of sentences

- \mathcal{M} **satisfies** Φ ($\mathcal{M} \models \Phi$), if for **every** sentence $A \in \Phi$, $\mathcal{M} \models A$
- If $\mathcal{M} \models \Phi$, say \mathcal{M} is a **model** of Φ
- Say that Φ is **satisfiable** if there is a structure \mathcal{M} such that $\mathcal{M} \models \Phi$

Unsatisfiable: if A is a **logical consequence** of Φ , then there is no \mathcal{M} such that $\mathcal{M} \models \Phi \cup \{\neg A\}$

Resolution by Refutation

Knowledge Base: a collection of **sentences** that represent what the agent believes about the world

Sentences in KB are explicit knowledge; logical consequences of the KB are implicit

Resolution works with formulas expressed in **clausal form**

Literal: an atomic formula or the negation of an a.f. (ex. $dog(Fido)$, $\neg cat(fido)$)

Clause: disjunction of literals (ex. $P(x) \vee \neg Q(x, y)$ $\neg O(fido) \vee \neg Dog(fido)$)

Clausal Theory: a set of clauses, can be considered as conjunction of clauses

$(P(x) \vee \neg Q(x, y), \neg O(fido))$

Resolution Proof using inference rule

$$\frac{a_1 \vee a_2 \vee \dots \vee a_n \vee c \quad b_1 \vee b_2 \vee \dots \vee b_m \vee \neg c}{a_1 \vee a_2 \vee \dots \vee a_n \vee b_1 \vee b_2 \vee \dots \vee b_m}$$

Resolution by Refutation to Show: **KMA**

- Assume **$\neg A$** is true to generate a contradiction (**refutation**)
- Convert $\neg A$ and all sentences in KM to a **clausal theory** C
- **Resolve** the clauses in C until an empty clause is obtained

Eliminate Implications

Implication Rule: $A \rightarrow B \iff \neg A \vee B$

- $\neg(A \wedge B) \iff \neg A \vee \neg B$
- $\neg(A \vee B) \iff \neg A \wedge \neg B$
- $\neg \forall x A \iff \exists x \neg A$
- $\neg \exists x A \iff \forall x \neg A$

Standardize Variables: rename variables so that each quantified variable is unique

Skolemizaation: remove existential quantifiers by introducing new function symbols

Convert to Prenex Form: bring all quantifiers to the front

(1) $\forall x P \wedge Q \iff Q \wedge \forall x P \iff \forall x (P \wedge Q)$; (2) $\forall x P \vee Q \iff Q \vee \forall x P \iff \forall x (P \vee Q)$;

Conjunctions over disjunctions: $A \vee (B \wedge C) \iff (A \vee B) \wedge (A \vee C)$

Flatten nested \wedge, \vee : $A \vee (B \vee C)$ to $(A \vee B \vee C)$

Convert to Clauses

Resolution is refutation complete: If a set of clauses is unsatisfiable (i.e., when the answer is "YES") and so some branch contains $[]$, a breadth-first search guaranteed to find $[]$.

First-order unsatisfiability is semi-decidable, but not decidable. Thus, calculating entailments is semi-decidable and undecidable. First-order satisfiability is undecidable.

Decidable if there is some algorithm that correctly generates a "YES-NO" answer for every possible input. Otherwise, it's undecidable.

Semi-decidable if there is some algorithm that correctly generates "YES" answers, but does not terminate on some inputs for which the answer is "NO"